



Cairo University - Faculty Of Engineering
Computer Engineering Department
Natural Language Processing - Fall 2025



Final Project Submission Document

Arabic Diacritization Pipelines

Submitted to
Eng. Omar Samir

Submitted by

Name	Section	BN	ID
Ahmed Hamed	1	3	9220027
Akram Hany	1	13	9220158
Amir Anwar	1	14	9220166
Somia Saad	1	26	9202666

December 10, 2025

Project Repository

The complete source code and trained model weights for this project can be found in our GitHub repository:

<https://github.com/El-Maktab/Arabic-Diacritization>

Pipeline Architectures

We implemented **two** main pipeline architectures:

1. Pipeline A — Character-Level Sequential Models
2. Pipeline B — Hierarchical Word + Character Model with Attention

1 Pipeline A — Character-Level Encoder + Diacritic Classifier

This pipeline treats the task as a sequence labeling problem at the character level.

1. Data Preprocessing

1.1 Data Cleaning We removed all characters outside the allowed vocabulary:

- Arabic letters
- Arabic diacritics
- Spaces
- Standard punctuation
- Padding symbol

Any invalid tokens (HTML tags, English characters, digits, unsupported Unicode symbols) were stripped.

1.2 Sentence Tokenization We split each sentence into smaller segments based on punctuation to reduce excessively long sequences and stabilize LSTM gradients.

1.3 Character/Diacritic Separation Each Arabic character was decomposed into:

- **Base letter**
- **Associated diacritic** (or PAD if none)

This enabled independent modeling of diacritics as categorical labels.

1.4 Encoding We performed:

- **Character encoding:** letter → ID
- **Diacritic encoding:** diacritic → ID
- **Sequence padding:** fixed-length representation for batching

2. Feature Extraction

We implemented and evaluated several feature families:

2.1 Trainable Embeddings (BiLSTM / BiLSTM-CRF / CNN-CRF)

- Learned jointly with the model during training
- Embedding dimension tuned experimentally (128–256)

2.2 Skip-gram Word Embeddings (CRF Model Only) Skip-gram was used in the pure CRF model to provide word-level distributional semantics. *Hyperparameters*:

- Embedding dim: 100
- Window size: 5
- Workers: 4
- Min count: 1
- Epochs: 10

2.3 CNN Character Encoder (for CNN-CRF model) We extracted n-gram-like morphological features using a shallow CNN:

- Kernel sizes: (2, 3, 4)
- Filters per kernel: configurable
- ReLU activation + max-pooling

2.4 HMM Features For the HMM model, we computed:

- **Initial probabilities**
- **Transition probabilities** between diacritics
- **Emission probabilities** $p(\text{letter}|\text{diacritic})$

3. Models

We trained four main models under Pipeline A.

3.1 Bidirectional LSTM

A standard sequence labeling model using only trainable embeddings.

Hyperparameters

- Embedding dim: 128
- Hidden dim: 256
- Layers: 3
- Dropout: 0.2
- Batch size: 32
- Epochs: 5
- Learning rate: 0.001

3.2 Conditional Random Field (CRF)

Uses skip-gram embeddings + handcrafted contextual features.

Skip-gram hyperparameters: Embedding dim: 100, Window size: 5, Min count: 1, SG = 1, Epochs: 10.

CRF hyperparameters:

- Algorithm: L-BFGS
- C1: 0.1 (L1 regularization)
- C2: 0.1 (L2 regularization)
- Max iterations: 100
- Context window: ± 2 characters

3.3 CNN-CRF Model

Combines CNN character encoder + skip-gram embeddings + CRF inference.

- CNN extracts morphological features
- CRF ensures valid label transitions

CNN hyperparameters: Character embedding dim: 30, Filters: 50, Kernel sizes: 2, 3, 4, CNN output dim: 150, Batch size: 512.

3.4 Bidirectional LSTM-CRF

This hybrid model adds a CRF decoding layer on top of the BiLSTM output, allowing explicit modeling of dependencies between consecutive diacritics.

Hyperparameters: Same as BiLSTM **except**: Batch size: 128, Epochs: 6. *This model ultimately achieved the best performance.*

3.5 Hidden Markov Model (HMM)

A baseline generative model trained using Initial probability estimation, Transition matrix, and Emission matrix. Inference performed using Viterbi decoding.

2 Pipeline B — Hierarchical Word + Character Model

4.1 Overview

Pipeline B implements the **D2** (**T**wo-**L**evel **D**iacritizer) architecture from the paper *"Effective Deep Learning Models for Automatic Diacritization of Arabic Text"*. Unlike Pipeline A which processes text character-by-character, this architecture separates word-level and character-level processing.

Key insight: Diacritization requires both local morphological features (within-word patterns) and global syntactic context (cross-word dependencies like grammatical agreement).

4.2 Preprocessing Differences

- **Sliding Window Segmentation:** Sentences are split into overlapping segments of $T_s = 10$ words with $\text{stride} = 5$.
- **Word-Level Tokenization:** Text is tokenized at the word level. Words below $\text{min_freq} = 2$ are mapped to `<UNK>`.
- **Character-Diacritic Separation:** Combined diacritics (e.g., shaddah + vowel) mapped to one of 15 classes.

4.3 Feature Extraction

- **Trainable Word Embeddings:** Dim 128.
- **Trainable Character Embeddings:** Dim 32.
- **Contextual Word Features (BiLSTM):** 2-layer BiLSTM, Hidden dim 256.
- **Contextual Character Features (BiLSTM):** Character embeddings concatenated with parent word context, processed by 3-layer BiLSTM (Hidden 512).
- **Cross-Level Attention Features:** Scaled dot-product attention over word encodings. For each character, attends to all other words.

4.4 Model Architecture

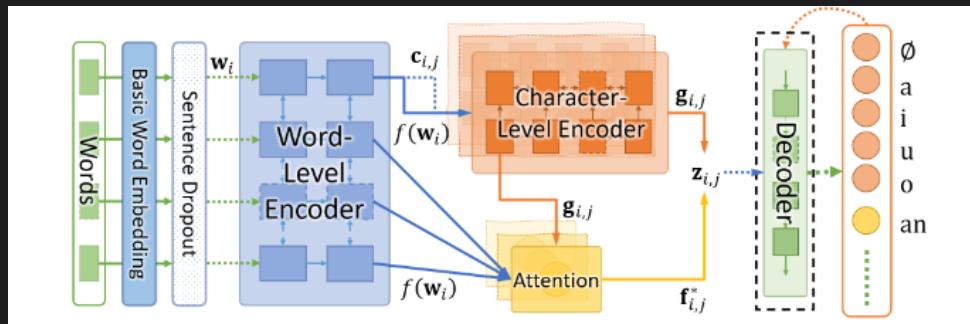


Figure 1: Architecture of the D2 Hierarchical Model

Table 1: Component Specification

Component	Specification
Word Encoder	2-layer BiLSTM, hidden=256, word_dropout=0.2
Char Encoder	3-layer BiLSTM, hidden=512, input=[char_emb; word_ctx]
Attention	Scaled dot-product, Q=char, K/V=words, self-masked
Classifier	Linear(1024 → 15 classes)

Table 2: Training Hyperparameters

Parameter	Value
Batch size	128
Learning rate	0.002
Optimizer	Adam
LR scheduler	ReduceLROnPlateau (factor=0.5, patience=1)
Dropout	0.25 (vertical), 0.2 (word/input)
Early stopping	3 epochs patience

4.5 D2 Experiments

Exp	Acc	CE Acc	W-L	C-L	W-Hid	C-Hid	Drop	W-Drop	Batch	LR	MaxW	Str	Notes
1	96.1	92.8	2	2	256	512	0.25	0.2	32	0.002	10	5	Baseline config
2	95.0	93.0	2	2	256	512	0.25	0.0	32	0.002	10	5	No word dropout
3	96.87	95.0	2	3	256	512	0.25	0.2	32	0.002	10	5	+1 char layer
4	81.0	-	2	5	256	512	0.25	0.2	32	0.002	10	5	Overfitted
5	96.98	94.77	2	3	256	512	0.25	0.2	128	0.002	20	15	Larger batch
6	96.97	94.81	2	4	256	512	0.25	0.2	128	0.002	20	15	+1 char layer
7	97.2	95.37	4	4	256	512	0.25	0.2	128	0.002	15	15	Best

Table 3: D2 Model Experiments (W-L: Word Layers, C-L: Char Layers, Str: Stride)

Best D2 Model Results:

- Overall Accuracy: **97.20%**
- Case Ending Accuracy: **95.37%**
- DER (Diacritic Error Rate): **2.80%**

5. Evaluation

Accuracy Evaluation

Table 4: Case Ending Accuracy

Model	Accuracy
BiLSTM	95.69
CRF	83.82
CNN-CRF	85.17
BiLSTM-CRF	95.75
HMM	45.81
D2 Model	95.37

Table 5: Overall Performance

Model	Accuracy
BiLSTM	97.73
CRF	82.33
CNN-CRF	83.67
BiLSTM-CRF	97.80
HMM	61.41
D2 Model	97.20

6. Final Model Selection

We selected the **Bidirectional LSTM-CRF** model as our final submission.

Reasons for selection:

- It achieved the **highest accuracy** among all tested models.
- It obtained the **lowest Diacritic Error Rate (DER)**.
- CRF decoding significantly reduced invalid or unlikely diacritic sequences.
- Better generalization to rare diacritics and complex morphological patterns.